

OVERVIEW

I am a final-year Ph.D. candidate in Statistics at The University of Texas at Austin. My research focuses on the intersection of optimal transport, statistics, and machine learning.

EDUCATION

The University of Texas at Austin Ph.D. in Statistics at Department of Statistics and Data Sciences	Texas, USA 2021–Present
<ul style="list-style-type: none">– Expected graduation date: May, 2026.– Advisors: Professor Nhat Ho and Professor Peter Mueller.	
Hanoi University of Science and Technology (HUST) B.Sc in Computer Science (5 years program)	Hanoi, Vietnam 2015–2020
<ul style="list-style-type: none">– Top: 1%, graduated with Excellent Degree.	

ACADEMIC EXPERIENCE

The University of Texas at Austin <i>Graduate Research Assistant</i>	Texas, USA January, 2025 –May, 2025
<ul style="list-style-type: none">– Research topic: Machine Learning for Human Epilepsy Analysis.	
The University of Texas at Austin <i>Graduate Teaching Assistant</i>	Texas, USA August, 2024 –January, 2025
<ul style="list-style-type: none">– Courses: Linear Models, Elements of Statistics.	
The University of Texas at Austin <i>Graduate Research Assistant</i>	Texas, USA September, 2023 –May, 2024
<ul style="list-style-type: none">– Research topic: Effective and Scalable Transportation Metrics for Machine Learning and Statistics.	
The University of Texas at Austin <i>Graduate Research Assistant</i>	Texas, USA September, 2022 –May, 2023
<ul style="list-style-type: none">– Research topic: Large-scale Optimal Transport for Machine Learning.	

INDUSTRIAL EXPERIENCE

Amazon Science <i>Applied Scientist Intern</i>	Austin, TX, USA May, 2025 –August, 2025
<ul style="list-style-type: none">– Developed a multi-task training framework session-based recommendations with large language models and improved online-inference by using KV caching and speculative decoding.	
Amazon Science <i>Applied Scientist Intern</i>	Seattle, WA, USA May, 2024 –August, 2024
<ul style="list-style-type: none">– Proposed a framework that leverages large language models and deep generative models to increase the diversity of product recommendations.	

Toyota InfoTech Labs*Research Intern*

Mountain View, CA, USA

May, 2023 – August, 2023

- Developed a Transformer-based framework with global-local decomposition, improving battery health prediction.

AT&T Labs*Research Intern*

Texas, USA

June, 2022 – August, 2022

- Proposed and implemented co-clustering algorithms to analyze user browsing behavior in PySpark.

VinAI Research (Acquired by Qualcomm AI Research)*AI Research Resident*

Hanoi, Vietnam

2019 – 2021

- Did research on Deep Generative Models (VAEs, GANs, score matching, diffusion models) and improved them with Optimal Transport (sliced Wasserstein distance, Sinkhorn divergence).

BOOKS

1. **K. Nguyen**, “An introduction to sliced optimal transport”, *Accepted and under final production at Foundations and Trends® in Computer Graphics and Vision*, 2025.
2. BDL Consortium, “Bayesian deep learning”, *CRC Press*, 2025, David Dunson, Vincent Fortuin, Yarin Gal, Philipp Hennig, Konstantina Palla, Theodore Papamarkou, Tim Rudner, Sara Wade, Andrew Wilson, Cheng Zhang, and **Khai Nguyen** among other contributing authors; under preparation.

UNDER REVIEW, UNDER REVISION, AND UNDER PREPARATION

(*) denotes equal contribution

1. **K. Nguyen***, H. Nguyen*, T. Pham, and N. Ho, “Fast estimation of Wasserstein distances via regression on sliced Wasserstein distances”, *Under review*, 2026.
2. **K. Nguyen**, Y. Ni, and P. Mueller, “Bayesian multiple multivariate density-density regression”, *To be submitted*, 2025.
3. **K. Nguyen**, “Streaming sliced optimal transport”, *To be submitted*, 2025.
4. **K. Nguyen**, Y. Ni, and P. Mueller, “Bayesian multivariate density-density regression”, *Under review*, 2025.
5. T. Nguyen M., N. Tran N., **K. Nguyen**, and R. Baraniuk, “Improving routing in sparse mixture of experts with graph of tokens”, *Under review*, 2025.
6. N. Barileto, **K. Nguyen**, and N. Ho, “Data-driven DRO and economic decision theory: An analytical synthesis with Bayesian nonparametric advancements”, *Under review*, 2024.
7. **K. Nguyen** and P. Mueller, “Summarizing Bayesian nonparametric mixture posterior - sliced optimal transport metrics for Gaussian mixtures”, *Under second revision at Journal of Computational and Graphical Statistics (JCGS)*, 2024.

REFEREED PUBLICATIONS

(*) denotes equal contribution,

Google Scholar: <https://scholar.google.com/citations?user=im5fNaQAAAAJ&hl=en>

1. M. Luong, **K. Nguyen**, D. Phung, G. Haffari, and L. Qu, “Unbiased sliced Wasserstein kernels for high-quality audio captioning”, *Neural Information Processing Systems (NeurIPS)*, 2025.
2. **K. Nguyen***, H. Nguyen*, T. Pham, and N. Ho, “Lightspeed geometric dataset distances via sliced optimal transport”, *International Conference on Machine Learning (ICML)*, 2025.
3. **K. Nguyen***, H. Nguyen*, and N. Ho, “Towards marginal fairness sliced Wasserstein barycenter”, *International Conference on Learning Representations (ICLR)*, **Spotlight 3.2%**, 2025.
4. **K. Nguyen** and N. Ho, “Hierarchical hybrid sliced Wasserstein: A scalable metric for heterogeneous joint distributions”, *Neural Information Processing Systems (NeurIPS)*, 2024.
5. **K. Nguyen**, S. Zhang, T. Le, and N. Ho, “Sliced Wasserstein with random-path projecting directions”, *International Conference on Machine Learning (ICML)*, 2024.
6. T. T. Le, **K. Nguyen**, S. Sun, N. Ho, and X. Xie, “Integrating efficient optimal transport and functional maps for unsupervised shape correspondence learning”, in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
7. **K. Nguyen**, N. Barileto, and N. Ho, “Quasi-Monte Carlo for 3D sliced Wasserstein”, *International Conference on Learning Representations (ICLR)*, **Spotlight 5%**, 2024.
8. **K. Nguyen** and N. Ho, “Sliced Wasserstein estimation with control variates”, *International Conference on Learning Representations (ICLR)*, 2024.
9. T. Le, **K. Nguyen**, N. Ho, S. Sun, K. Han, and X. Xie, “Diffeomorphic deformation via sliced Wasserstein distance optimization for cortical surface reconstruction”, *International Conference on Learning Representations (ICLR)*, 2024.
10. M. Luong, **K. Nguyen**, N. Ho, R. Haf, D. Phung, and L. Qu, “Revisiting deep audio-text retrieval through the lens of transportation”, *International Conference on Learning Representations (ICLR)*, 2024.
11. H. Nguyen, **K. Nguyen**, and N. Ho, “On parameter estimation in deviated Gaussian mixture of experts”, *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2024.
12. H. Nguyen, T. Nguyen, **K. Nguyen**, and N. Ho, “Towards convergence rates for parameter estimation in Gaussian-gated mixture of experts”, *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2024.
13. **K. Nguyen** and N. Ho, “Energy-based sliced Wasserstein distance”, *Neural Information Processing Systems (NeurIPS)*, 2023.
14. **K. Nguyen**, T. Ren, and N. Ho, “Markovian sliced Wasserstein distances: Beyond independent projections”, *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
15. D. Le*, H. Nguyen*, **K. Nguyen***, T. Nguyen, and N. Ho, “Fast approximation of the generalized sliced-Wasserstein distance”, *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2024.
16. X. Han, T. Ren, T. M. Nguyen, **K. Nguyen**, J. Ghosh, and N. Ho, “Robustify Transformers with robust kernel density estimation”, *Neural Information Processing Systems (NeurIPS)*, 2023.
17. D. Do, H. Nguyen, **K. Nguyen**, and N. Ho, “Minimax optimal rate for parameter estimation in multivariate deviated models”, *Neural Information Processing Systems (NeurIPS)*, 2023.

18. **K. Nguyen***, D. Nguyen*, and N. Ho, “Self-attention amortized distributional projection optimization for sliced Wasserstein point-clouds reconstruction”, *International Conference on Machine Learning (ICML)*, 2023.
19. **K. Nguyen**, T. Ren, H. Nguyen, L. Rout, T. Nguyen, and N. Ho, “Hierarchical sliced Wasserstein distance”, *International Conference on Learning Representations (ICLR)*, 2023.
20. D. Nguyen, T. Nguyen, **K. Nguyen**, D. Phung, H. Bui, and N. Ho, “On cross-layer alignment for model fusion of heterogeneous neural networks”, *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023.
21. **K. Nguyen** and N. Ho, “Revisiting sliced Wasserstein on images: From vectorization to convolution”, *Neural Information Processing Systems (NeurIPS)*, 2022.
22. **K. Nguyen** and N. Ho, “Amortized projection optimization for sliced Wasserstein generative models”, *Neural Information Processing Systems (NeurIPS)*, 2022.
23. T. Nguyen, M. Pham, T. Nguyen, **K. Nguyen**, S. J. Osher, and N. Ho, “Transformer with Fourier integral attentions”, *Neural Information Processing Systems (NeurIPS)*, 2022.
24. T. Nguyen, T. Nguyen, H. Do, **K. Nguyen**, V. Saragadam, M. Pham, K. Nguyen, N. Ho, and S. J. Osher, “Improving transformer with an admixture of attention heads”, *Neural Information Processing Systems (NeurIPS)*, 2022.
25. **K. Nguyen***, D. Nguyen*, T.-A. Vu-Le, T. Pham, and N. Ho, “Improving mini-batch optimal transport via partial transportation”, in *International Conference on Machine Learning (ICML)*, 2022.
26. **K. Nguyen**, D. Nguyen, Q. Nguyen, T. Pham, H. Bui, D. Phung, T. Le, and N. Ho, “On transportation of mini-batches: A hierarchical approach”, in *International Conference on Machine Learning (ICML)*, 2022.
27. K. Le, H. Nguyen, **K. Nguyen**, T. Pham, and N. Ho, “On multimarginal partial optimal transport: Equivalent forms and computational complexity”, in *International Conference on Artificial Intelligence and Statistics (AISTATS)*, PMLR, 2022, pp. 4397–4413.
28. S. Nguyen, D. Nguyen, **K. Nguyen**, K. Than, H. Bui, and N. Ho, “Structured dropout variational inference for bayesian neural networks”, *Neural Information Processing Systems (NeurIPS)*, 2021.
29. **K. Nguyen**, S. Nguyen, N. Ho, T. Pham, and H. Bui, “Improving relational regularized autoencoders with spherical sliced fused Gromov-Wasserstein”, in *International Conference on Learning Representations (ICLR)*, 2021.
30. **K. Nguyen**, N. Ho, T. Pham, and H. Bui, “Distributional sliced-Wasserstein and applications to generative modeling”, in *International Conference on Learning Representations (ICLR)*, **Spotlight 3.78%**, 2021.

AWARDS

-
- **BNP14 Travel Award.** 2025
 - **UT Austin Graduate Continuing Fellowship.** 2025
The UT Austin Graduate Continuing Fellowship is one of the university’s most prestigious awards, recognizing exceptional academic achievement and research excellence among continuing graduate students.
 - **Top Reviewer Award at NeurIPS.** 2024
 - **ICML Travel Grants.** 2023

- **Top Reviewer Award at NeurIPS.** 2022
- **NeurIPS Scholar Award.** 2022
- **ICML Travel Grants.** 2022
- **UT Austin Graduate Recruitment Fellowship.** 2021
The UT Austin Graduate Recruitment Fellowship is a highly prestigious award offered to attract top academic talents.

TALKS

- The Bayesian Young Statisticians Meeting (Talk in session with discussion): “Summarizing Bayesian Nonparametric Mixture Posterior - Sliced Optimal Transport Metrics for Gaussian Mixtures”, 2025.
- International Conference on Bayesian Nonparametrics (Contributed Talk): “Summarizing Bayesian Nonparametric Mixture Posterior - Sliced Optimal Transport Metrics for Gaussian Mixtures”, 2025.
- International Conference on Machine Learning (ICML) (Spotlight Talk): “On Transportation of Mini-batches: A Hierarchical Approach”, 2022.
- International Conference on Machine Learning (ICML) (Spotlight Talk): “Improving Mini-batch Optimal Transport via Partial Transportation”, 2022.

PROFESSIONAL SERVICES

- Reviewer at Journal of Machine Learning Research (JMLR).
- Reviewer at Statistics and Computing.
- Reviewer at IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI).
- Reviewer at Transactions on Machine Learning Research (TMLR).
- Reviewer at IEEE Transactions on Information Theory.
- Reviewer at Machine Learning Journal.
- Reviewer at International Conference on Machine Learning (ICML) 2022-2025.
- Reviewer at Workshop on Challenges in Deployable Generative AI (ICML) 2023.
- Reviewer at Conference on Neural Information Processing Systems (NeurIPS) 2021-2025.
- Reviewer at Workshop on Deep Generative Models (NeurIPS) 2021.
- Reviewer at International Conference on Learning Representations (ICLR) 2022-2026.
- Reviewer at International Conference on Artificial Intelligence and Statistics (AISTATS) 2022-2025.
- Reviewer at AAAI Conference on Artificial Intelligence (AAAI) 2023-2026 .
- Reviewer at IEEE / CVF Computer Vision and Pattern Recognition Conference (CVPR) 2023-2026.
- Reviewer at International Conference on Computer Vision (ICCV) 2023.
- Reviewer at European Conference on Computer Vision (ECCV) 2024.
- Reviewer at IEEE/CVF Winter Conference on Applications of Computer Vision (WACV) 2026.
- Reviewer at IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP) 2024-2025.
- Reviewer at Conference on Language Modeling (COLM) 2024-2025.
- Reviewer at Learning on Graphs Conference (LOG) 2024-2025.

MENTORSHIP

- Hai Nguyen (Undergraduate student at Hanoi University of Science and Technology).

PROFESSIONAL MEMBERSHIP

- Institute of Electrical and Electronics Engineers (IEEE) *Regular Member*.
- International Society for Bayesian Analysis (ISBA) *Student Member*.
- Institute of Mathematical Statistics (IMS) *Student Member*.

TECHNICAL SKILLS

- **Python:** Proficient.
Libraries: Pytorch (proficient), Scikit-Learn (proficient), Numpy (proficient), Pandas (proficient), Matplotlib (proficient), Pyspark (basic), and so on.
- **Developer Tools:** Git.
- **Systems:** Linux.

REFERENCES

- **Peter Mueller.** Email: pmueller@math.utexas.edu (PhD advisor).
- **Nhat Ho.** Email: minhnhhat@utexas.edu (PhD advisor).
- **Yang Ni.** Email: yni@stat.tamu.edu (collaborator).